

Chapter 4

Nonparametric Techniques



Bayes Theorem for Classification

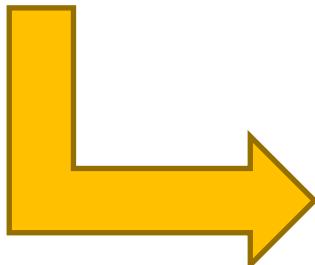
$$P(\omega_j | \mathbf{x}) = \frac{p(\mathbf{x} | \omega_j) \cdot P(\omega_j)}{p(\mathbf{x})} \quad (1 \leq j \leq c) \quad (\text{Bayes Formula})$$

To compute posterior probability $P(\omega_j | \mathbf{x})$, we need to know:

Prior probability: $P(\omega_j)$

Likelihood: $p(\mathbf{x} | \omega_j)$

- Case I: $p(\mathbf{x} | \omega_j)$ has certain **parametric form** $p(\mathbf{x} | \omega_j, \theta_j)$



Maximum-Likelihood (ML) Estimation

Bayesian Parameter Estimation

Bayes Theorem for Classification (Cont.)

Potential problems for Case I

The assumed parametric form **may not fit the ground-truth density** encountered in practice, e.g.:

Assumed parametric form: Unimodal (单峰, such as Gaussian pdf)

Ground-truth form: Multimodal (多峰)

□ **Case II:** $p(\mathbf{x}|\omega_j)$ doesn't have **parametric form**

Let the data
speak for
themselves!



Parzen Windows

k_n -nearest-neighbor

Density Estimation

General settings

Feature space: $\mathcal{F} = \mathbf{R}^d$

Feature vector: $\mathbf{x} \in \mathcal{F}$

pdf function: $\mathbf{x} \sim p(\cdot)$



How to estimate
 $p(\mathbf{x})$ from the
training examples?

Fundamental fact

The probability of a vector \mathbf{x} **falling into a region** $\mathcal{R} \subset \mathcal{F}$:

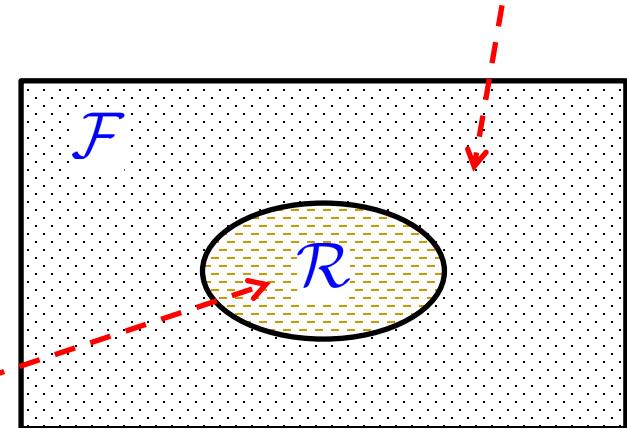
$$P = \Pr[\mathbf{x} \in \mathcal{R}] = \int_{\mathcal{R}} p(\mathbf{x}') d\mathbf{x}'$$

A smoothed/averaged
version of $p(\mathbf{x})$

Density Estimation (Cont.)

$$\Pr[\mathbf{x} \notin \mathcal{R}] = 1 - P$$

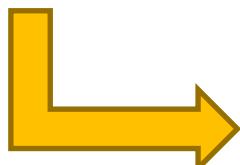
$$P = \Pr[\mathbf{x} \in \mathcal{R}] = \int_{\mathcal{R}} p(\mathbf{x}') d\mathbf{x}'$$



$$\Pr[\mathbf{x} \in \mathcal{R}] = P$$

Given n examples (i.i.d.) $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ with $\mathbf{x}_i \sim p(\cdot)$ ($1 \leq i \leq n$)

Let X be the (discrete) **random variable** representing the number of examples falling into \mathcal{R}



X will take Binomial distribution (二项分布):

$$X \sim \mathcal{B}(n, P)$$

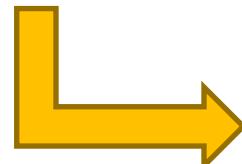
Density Estimation (Cont.)

$$\Pr[X = r] = \binom{n}{r} P^r (1 - P)^{n-r} \quad (0 \leq r \leq n)$$

$X \sim \mathcal{B}(n, P)$

$\mathcal{E}[X] = nP$

Table 3.1 [pp.109]


$$P = \frac{\mathcal{E}[X]}{n}$$

Assume \mathcal{R} is small

$$P = \Pr[\mathbf{x} \in \mathcal{R}] = \int_{\mathcal{R}} p(\mathbf{x}') d\mathbf{x}'$$

$p(\cdot)$ hardly varies
within \mathcal{R}

$$\simeq p(\mathbf{x}) \int_{\mathcal{R}} 1 d\mathbf{x}' \quad (\mathbf{x} \text{ is a point within } \mathcal{R})$$

$$P \simeq p(\mathbf{x}) V \quad (V \text{ is the volume enclosed by } \mathcal{R})$$



Density Estimation (Cont.)

$$\left. \begin{array}{l} P = \frac{\mathcal{E}[X]}{n} \\ P \simeq p(\mathbf{x}) V \end{array} \right\} p(\mathbf{x}) = \frac{\mathcal{E}[X]/n}{V} \quad \xrightarrow{\hspace{10em}} \quad p(\mathbf{x}) = \frac{k/n}{V}$$

$X \sim \mathcal{B}(n, P)$ **X peaks sharply
about $\mathcal{E}[X]$ when
 n is large enough**

Let k be the **actual value of X**
after observing the *i.i.d.* training
examples $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$

$$k \simeq \mathcal{E}[X]$$



Density Estimation (Cont.)

To show the explicit relationships with n :

$$\mathcal{R} \xrightarrow{\hspace{1cm}} \mathcal{R}_n \text{ (containing } \mathbf{x})$$
$$p(\mathbf{x}) = \frac{k/n}{V} \xrightarrow{\hspace{1cm}} p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

V_n : volume of \mathcal{R}_n n : # training examples

Quantities:

k_n : # training examples falling within \mathcal{R}_n

Fix V_n and determine k_n $\xrightarrow{\hspace{1cm}}$ Parzen Windows

Fix k_n and determine V_n $\xrightarrow{\hspace{1cm}}$ k_n -nearest-neighbor

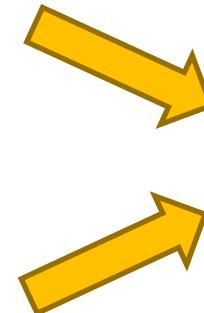


Parzen Windows

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

Fix V_n , and then determine k_n

Assume \mathcal{R}_n is a d -dimensional hypercube (超立方体)



The length of each edge is h_n

$$V_n = h_n^d$$



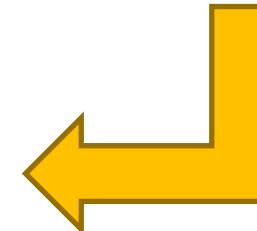
Determine k_n with **window function** (窗口函数),
a.k.a. **kernel function** (核函数), **potential function** (势函数), etc.

Emanuel Parzen
(1929-)

Parzen Windows (Cont.)

Window function: $\varphi(\mathbf{u}) = \begin{cases} 1 & |u_j| \leq 1/2; \quad j = 1, \dots, d \\ 0 & \text{otherwise} \end{cases}$

$\varphi(\mathbf{u})$ defines a **unit hypercube**
centered at the origin



$$\varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) = 1$$



\mathbf{x}_i falls within the hypercube
of volume V_n centered at \mathbf{x}



$$k_n = \sum_{i=1}^n \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right)$$

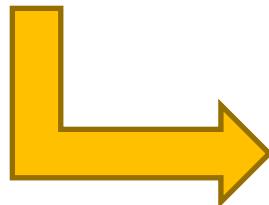
Parzen Windows (Cont.)

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n} \quad \longrightarrow \quad p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right)$$

$$k_n = \sum_{i=1}^n \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) \quad \longrightarrow$$

An average of functions
of \mathbf{x} and \mathbf{x}_i

$\varphi(\cdot)$ is not limited to be the hypercube window function of
Eq.9 [pp.164]



$\varphi(\cdot)$ could be any
pdf function:

$$\varphi(\mathbf{u}) \geq 0$$

$$\int \varphi(\mathbf{u}) d\mathbf{u} = 1$$

Parzen Windows (Cont.)

$$p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) \quad (V_n = h_n^d)$$

$\varphi(\cdot)$ being a pdf function  $p_n(\cdot)$ being a pdf function

$$\int p_n(\mathbf{x}) d\mathbf{x} = \frac{1}{nV_n} \sum_{i=1}^n \int \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) d\mathbf{x}$$

Integration by substitution (换元积分)

$$= \frac{1}{nV_n} \sum_{i=1}^n \int h_n^d \varphi(\mathbf{u}) d(\mathbf{u}) = \frac{1}{n} \sum_{i=1}^n \int \varphi(\mathbf{u}) d(\mathbf{u}) = 1$$

Let $\mathbf{u} = (\mathbf{x} - \mathbf{x}_i)/h_n$

window function
(being pdf) $\varphi(\cdot)$ + window width h_n + training data \mathbf{x}_i  Parzen pdf $p_n(\cdot)$



Parzen Windows (Cont.)

Parzen pdf: $p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) \quad (V_n = h_n^d)$

$\varphi(\cdot)$ being a pdf function  $p_n(\cdot)$ being a pdf function

$$\delta_n(\mathbf{x}) = \frac{1}{V_n} \varphi\left(\frac{\mathbf{x}}{h_n}\right) \quad \Rightarrow \quad p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \delta_n(\mathbf{x} - \mathbf{x}_i)$$



What is the effect of
 h_n ("window width") on
the Parzen pdf?

- $p_n(\mathbf{x})$: **superposition** (叠加)
of n interpolations (插值)
- \mathbf{x}_i : contributes to $p_n(\mathbf{x})$ based
on its "**distance**" from \mathbf{x}
(i.e. " $\mathbf{x}-\mathbf{x}_i$ ")

Parzen Windows (Cont.)

The effect of h_n ("window width")

$$\delta_n(\mathbf{x}) = \frac{1}{V_n} \varphi\left(\frac{\mathbf{x}}{h_n}\right) = \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right)$$

Affects the *amplitude*
(vertical scale, 幅度)

What do "amplitude"
and "width" mean
for a function?

Affects the *width*
(horizontal scale, 宽度)

For $\varphi(\mathbf{u})$:

$$|\varphi(\mathbf{u})| \leq a \text{ (*amplitude*)}$$

$$|u_j| \leq b_j \text{ (*width*)} \quad (j = 1, \dots, d)$$

For $\delta_n(\mathbf{x})$:

$$|\delta_n(\mathbf{x})| \leq (1/h_n^d) \cdot a$$

$$|x_j| \leq h_n \cdot b_j \quad (j = 1, \dots, d)$$



Parzen Windows (Cont.)

$$\delta_n(\mathbf{x}) = \frac{1}{V_n} \varphi\left(\frac{\mathbf{x}}{h_n}\right) = \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right)$$

$\delta_n(\cdot)$ being a
pdf function

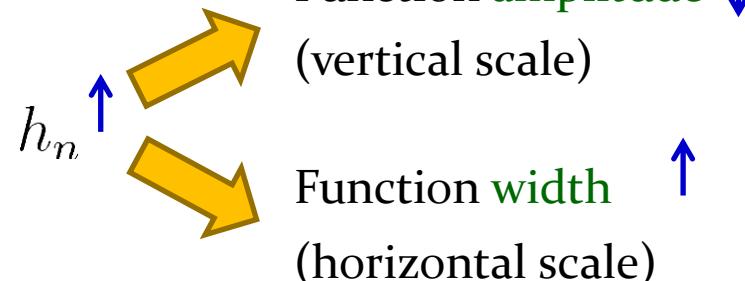
$$\int \delta_n(\mathbf{x}) d\mathbf{x} = \int \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right) d\mathbf{x}$$

Integration by substitution

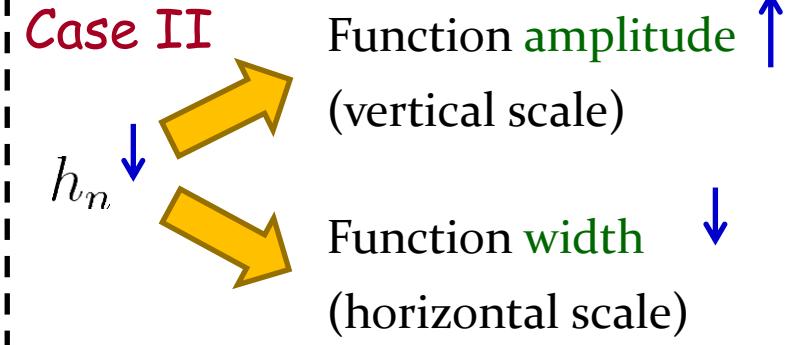
$$\text{Let } \mathbf{u} = \mathbf{x}/h_n$$

$$= \int \frac{1}{h_n^d} \cdot \varphi(\mathbf{u}) \cdot h_n^d d\mathbf{u} = \int \varphi(\mathbf{u}) d\mathbf{u} = 1$$

Case I



Case II

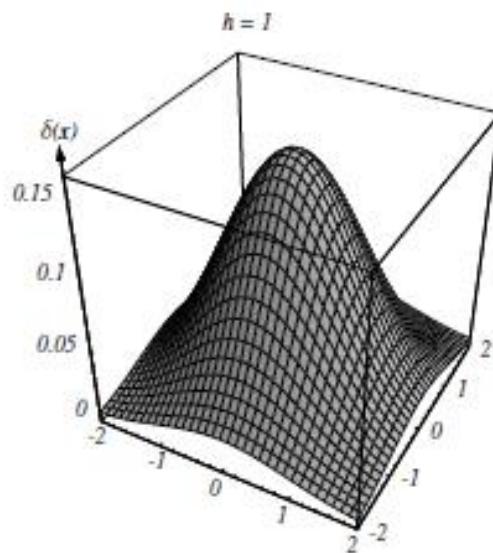


Parzen Windows (Cont.)

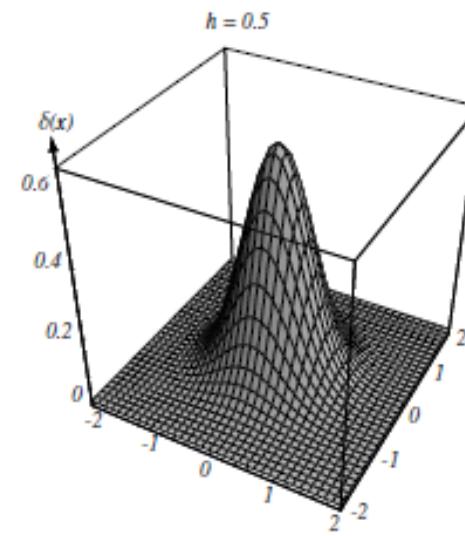
$$\delta_n(\mathbf{x}) = \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right)$$

Suppose $\varphi(\cdot)$ being a 2-d Gaussian pdf

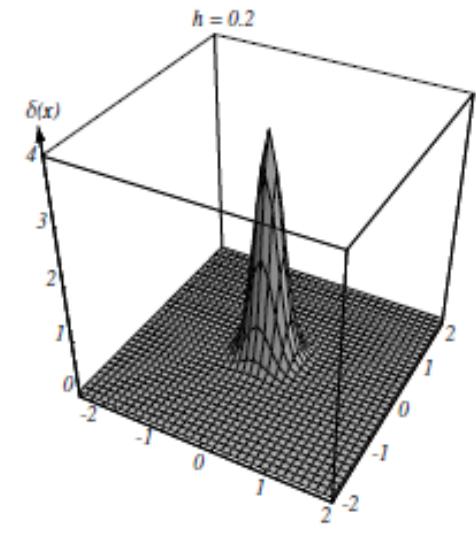
The shape of $\delta_n(\mathbf{x})$ with decreasing values of h_n



$h=1.0$



$h=0.5$



$h=0.2$

Parzen Windows (Cont.)

$$p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \delta_n(\mathbf{x} - \mathbf{x}_i), \text{ where } \delta_n(\mathbf{x}) = \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right)$$

□ h_n very large → $\delta_n(\mathbf{x})$ being *broad* with *small amplitude*

$p_n(\mathbf{x})$ will be the superposition of n broad, slowly changing (慢变) functions, i.e. being *smooth* (平滑) with *low resolution* (低分辨率)

□ h_n very small → $\delta_n(\mathbf{x})$ being *sharp* with *large amplitude*

$p_n(\mathbf{x})$ will be the superposition of n sharp pulses (尖脉冲), i.e. being *variable/unstable* (易变) with *high resolution* (高分辨率)



A *compromised value* (折衷值) of h_n should be sought for *limited number* of training examples

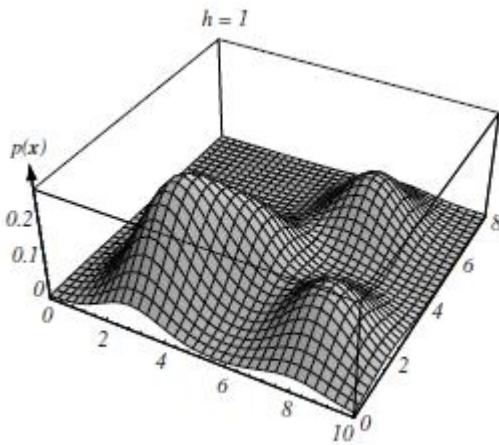
Parzen Windows (Cont.)

More illustrations:
Subsection 4.3.3 [pp.168]

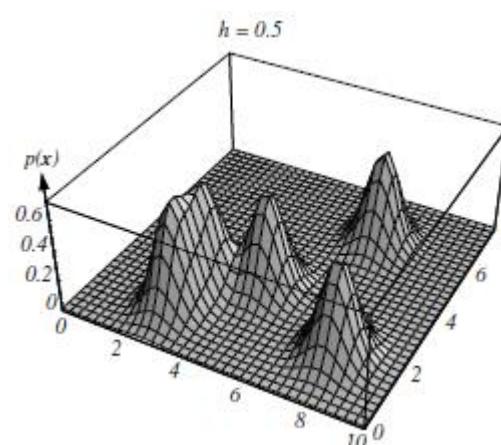
$$p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \delta_n(\mathbf{x} - \mathbf{x}_i), \text{ where } \delta_n(\mathbf{x}) = \frac{1}{h_n^d} \varphi\left(\frac{\mathbf{x}}{h_n}\right)$$

Suppose $\varphi(\cdot)$ being a 2-d Gaussian pdf and $n=5$

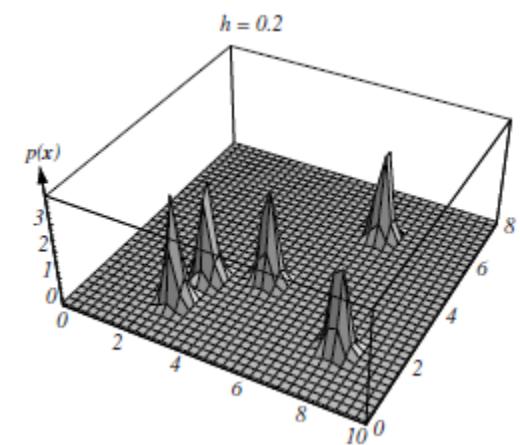
The shape of $p_n(\mathbf{x})$ with decreasing values of h_n



$h=1.0$



$h=0.5$



$h=0.2$

k_n -Nearest-Neighbor

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n}$$

Fix k_n , and then determine V_n

specify $k_n \rightarrow$ center a cell about $\mathbf{x} \rightarrow$ grow the cell until capturing k_n nearest examples \rightarrow return cell volume as V_n

The principled rule to specify k_n [pp.175]

$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$



A rule-of-thumb
choice for k_n :

$$k_n = \sqrt{n}$$

k_n -Nearest-Neighbor (Cont.)

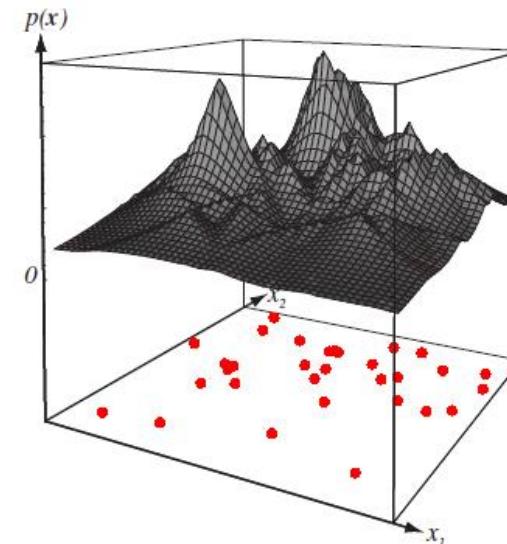
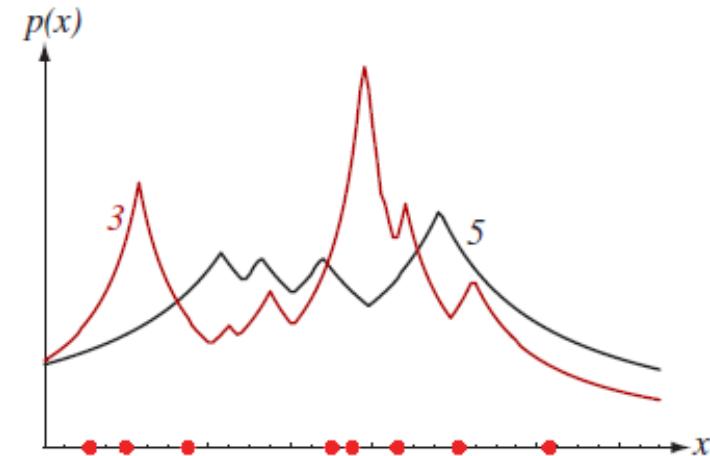
Eight points in one dimension
($n=8, d=1$)

red curve: $k_n=3$

black curve: $k_n=5$

Thirty-one points in two dimensions ($n=31, d=2$)

black surface: $k_n=5$



Related Topic

Nearest Neighbor Rule &
Distance Metric



Nearest-Neighbor (NN) Rule (最近邻准则)

Classification with nearest-neighbor rule

Given the label space $\Omega = \{\omega_1, \omega_2, \dots, \omega_c\}$ and a set of n labeled training examples $\mathcal{D}^n = \{(\mathbf{x}_i, \theta_i) \mid 1 \leq i \leq n\}$, where $\mathbf{x}_i \in \mathbf{R}^d$ and $\theta_i \in \Omega$

for test example \mathbf{x} , identify $\mathbf{x}' = \operatorname{argmin}_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} D(\mathbf{x}_i, \mathbf{x})$ and then assign the label θ' associated with \mathbf{x}' to \mathbf{x}

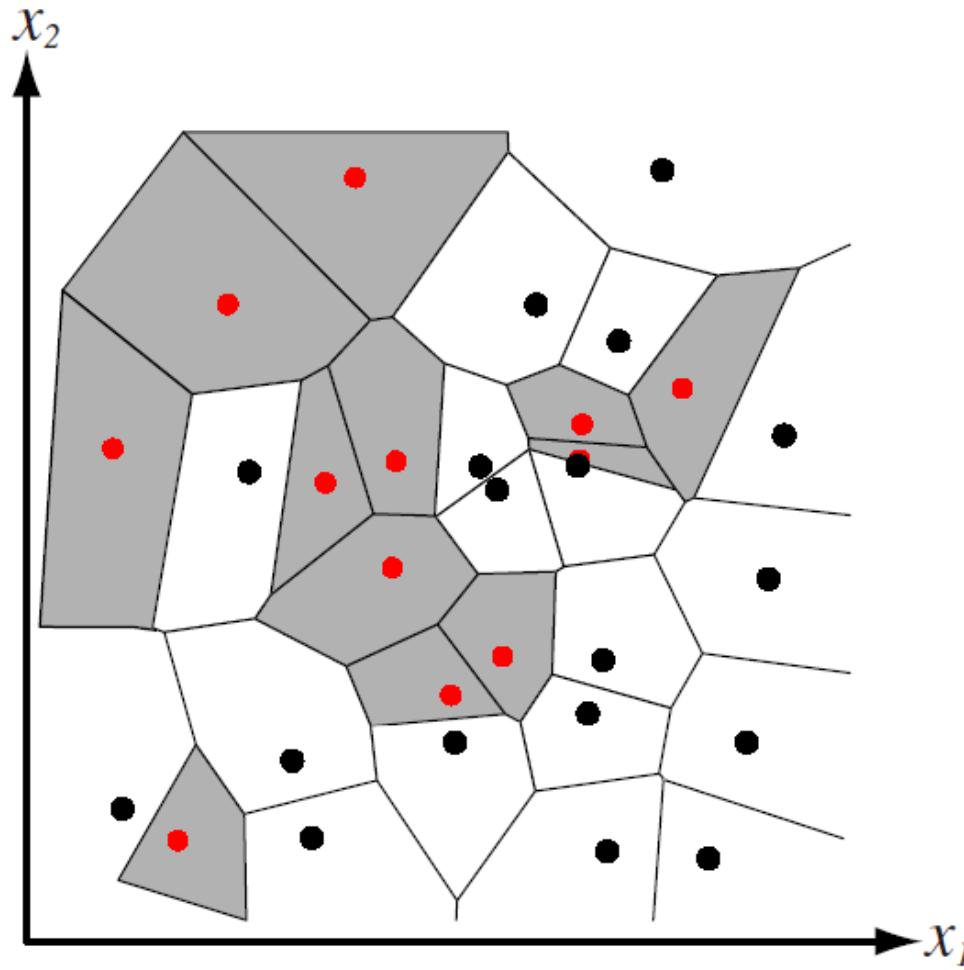
$D(\mathbf{a}, \mathbf{b})$: *distance metric between two vectors \mathbf{a} and \mathbf{b} , e.g. the Euclidean distance*

Basic assumption:

$$P(\omega_i \mid \mathbf{x}') \simeq P(\omega_i \mid \mathbf{x}) \\ \text{as } n \rightarrow \infty$$



Voronoi tessellation (维诺图)



Each training example \mathbf{x} leads to a cell in the Voronoi tessellation

- *any point in the cell is closer to \mathbf{x} than to any other training examples*
- *partition the feature space into n cells*
- *any point in the cell shares the same class label as \mathbf{x}*

Error Rate of Nearest Neighbor Rule

$P(e \mid \mathbf{x})$: The probability of making an erroneous classification on \mathbf{x} based on nearest-neighbor rule

$P(e)$: The **average probability of error** based on nearest-neighbor rule: $P(e) = \int P(e \mid \mathbf{x})p(\mathbf{x})d\mathbf{x}$

$P^*(e \mid \mathbf{x})$: The **minimum** possible value of $P(e \mid \mathbf{x})$, i.e. the one given by *Bayesian decision rule*: $P^*(e \mid \mathbf{x}) = 1 - \max_{1 \leq i \leq c} P(\omega_i \mid \mathbf{x})$

$P^*(e)$: The **Bayes risk** (*under zero-one loss*): $P^*(e) = \int P^*(e \mid \mathbf{x})p(\mathbf{x})d\mathbf{x}$

Error bounds of nearest neighbor rule

$$P^*(e) \leq P(e) \leq P^*(e) \left(2 - \frac{c}{c-1} P^*(e) \right) \quad (c: \# \text{ class labels})$$



k -Nearest-Neighbor (k NN) Rule (k -近邻准则)

Classification with k -nearest-neighbor rule

Given the label space $\Omega = \{\omega_1, \omega_2, \dots, \omega_c\}$ and a set of n labeled training examples $\mathcal{D}^n = \{(\mathbf{x}_i, \theta_i) \mid 1 \leq i \leq n\}$, where $\mathbf{x}_i \in \mathbf{R}^d$ and $\theta_i \in \Omega$

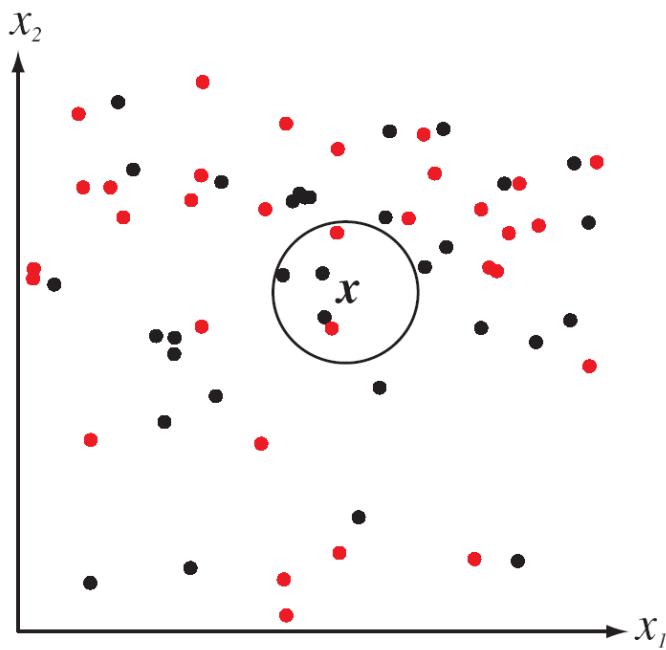
for test example \mathbf{x} , identify $S' = \{\mathbf{x}_i \mid \mathbf{x}_i \text{ is among the } k\text{NN of } \mathbf{x}\}$ and then assign the most frequent label w.r.t. S' , i.e. $\arg \max_{\omega_i \in \Omega} \sum_{\mathbf{x}_i \in S'} 1_{\theta_i = \omega_i}$ to \mathbf{x} .

1_π : an indicator function which returns 1 if predicate π holds, and 0 otherwise

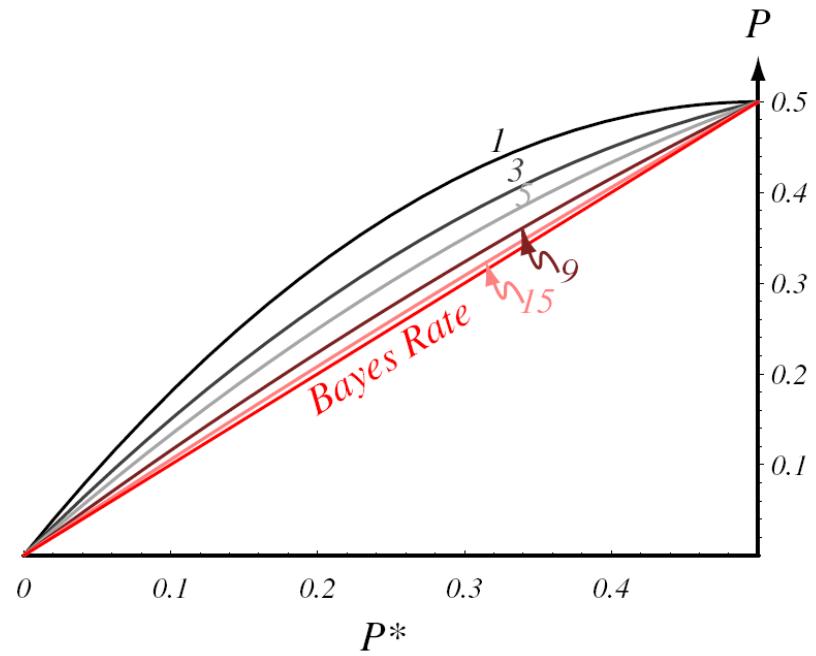
For binary classification problem ($c=2$), an odd value of k is generally used to avoid ties



k -Nearest-Neighbor (k NN) Rule (Cont.)



$c=2, k=5$



the error rate of k NN rule
(i.e. P) upper-bounded by
the Bayes risk (i.e. P^*) for
binary classification ($c=2$)

Computational Complexity of k NN Rule

Given n labeled training examples in d -dimensional feature space, the computational complexity of classifying one test example is $O(dn)$

General ways of reducing computational burden

□ **Partial distance:** $D_r(\mathbf{a}, \mathbf{b}) = \left(\sum_{j=1}^r (a_j - b_j)^2 \right)^{\frac{1}{2}} \quad (r < d)$

□ **Pre-structuring:**

*create some form of **search tree**, where nearest neighbors are recursively identified following the tree structure*

□ **Editing/Pruning/Condensing:**

*eliminate “**redundant**” (“useless”) examples from the training set, e.g. example surrounded by training examples of the same class label*

Properties of Distance Metric

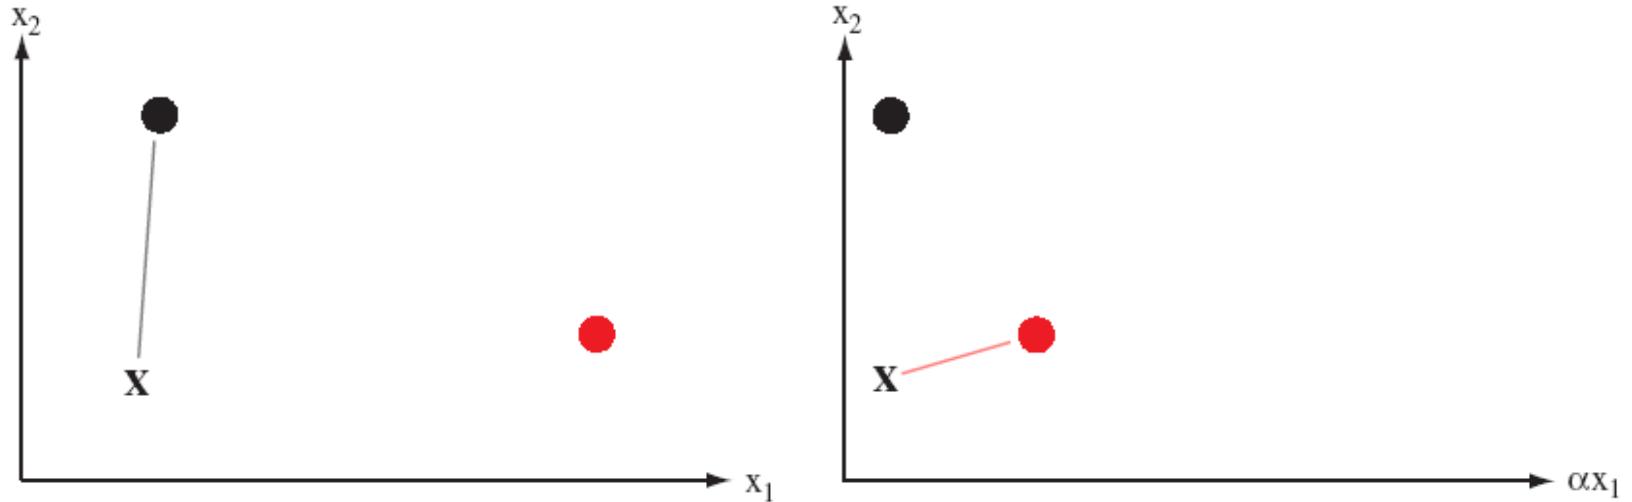
The NN/kNN rule depends on the use of distance metric to identify nearest neighbors

Four properties of distance metric

- **non-negativity:** $D(\mathbf{a}, \mathbf{b}) \geq 0$
- **reflexivity:** $D(\mathbf{a}, \mathbf{b}) = 0$ if and only if $\mathbf{a} = \mathbf{b}$
- **symmetry:** $D(\mathbf{a}, \mathbf{b}) = D(\mathbf{b}, \mathbf{a})$
- **triangle inequality:** $D(\mathbf{a}, \mathbf{b}) + D(\mathbf{b}, \mathbf{c}) \geq D(\mathbf{a}, \mathbf{c})$

Potential Issue of Euclidean Distance

$$D(\mathbf{a}, \mathbf{b}) = \left(\sum_{j=1}^d (a_j - b_j)^2 \right)^{\frac{1}{2}} \quad (\text{Euclidean distance})$$



Scaling the features → change the distance relationship

Possible solution: normalize each feature into equal-sized intervals, e.g. [0, 1]

Minkowski Distance Metric

$$L_k(\mathbf{a}, \mathbf{b}) = \left(\sum_{j=1}^d |a_j - b_j|^k \right)^{\frac{1}{k}} \quad (k > 0)$$

(a.k.a. L_k norm)

□ $k=2$: *Euclidean* distance

□ $k=1$: *Manhattan* distance (*city block* distance)

$$L_1(\mathbf{a}, \mathbf{b}) = \sum_{j=1}^d |a_j - b_j|$$

□ $k=\infty$: L_∞ distance

$$L_\infty(\mathbf{a}, \mathbf{b}) = \max_{1 \leq j \leq d} |a_j - b_j|$$



Distance Metric Between Sets

Tanimoto distance

$$D_{Tanimoto}(S_1, S_2) = \frac{n_1 + n_2 - 2n_{12}}{n_1 + n_2 - n_{12}}$$

$$(n_1 = |S_1|, n_2 = |S_2|, n_{12} = |S_1 \cap S_2|)$$

Example: treat each word as a set of characters

Which word out of ‘cat’, ‘pots’ and ‘patches’ mostly resembles ‘pat’?

cat

$$S_1 = \{p, a, t\}$$

$$D_{Tanimoto}(S_1, S_2) = \frac{3 + 3 - 2 * 2}{3 + 3 - 2} = 0.5$$

$$S_2 = \{c, a, t\}$$



$$D_{Tanimoto}(S_1, S_3) = \frac{3 + 4 - 2 * 2}{3 + 4 - 2} = 0.6$$

$$S_3 = \{p, o, t, s\}$$

$$D_{Tanimoto}(S_1, S_4) = \frac{3 + 7 - 2 * 3}{3 + 7 - 3} = 0.571$$

$$S_4 = \{p, a, t, c, h, e, s\}$$

Distance Metric Between Sets (Cont.)

Hausdorff distance

$$D_H(S_1, S_2) = \max \left(\max_{\mathbf{s}_1 \in S_1} \min_{\mathbf{s}_2 \in S_2} D(\mathbf{s}_1, \mathbf{s}_2), \max_{\mathbf{s}_2 \in S_2} \min_{\mathbf{s}_1 \in S_1} D(\mathbf{s}_2, \mathbf{s}_1) \right)$$

($D(\mathbf{s}_1, \mathbf{s}_2)$: any distance metric between \mathbf{s}_1 and \mathbf{s}_2)

Example: Hausdorff distance between two sets of feature vectors

$$S_1 = \{(0.1, 0.2)^t, (0.3, 0.8)^t\} \quad S_2 = \{(0.5, 0.5)^t, (0.7, 0.3)^t\}$$

$$\begin{aligned} D_H(S_1, S_2) &= \max (\max(0.5, 0.36), \max(0.36, 0.61)) \\ &= \max (0.5, 0.61) \\ &= 0.61 \end{aligned}$$



Summary

- Basic settings for nonparametric techniques
 - Let the data speak for themselves
 - Parametric form not assumed for class-conditional pdf
 - Estimate class-conditional pdf from training examples
 - Make predictions based on Bayes Formula
- Fundamental result in density estimation

n: # training examples

$$p_n(\mathbf{x}) = \frac{k_n/n}{V_n} \quad V_n: \text{volume of region } \mathcal{R}_n \text{ containing } \mathbf{x}$$

k_n: # training examples falling within \mathcal{R}_n



Summary (Cont.)

- Parzen Windows: Fix $V_n \rightarrow$ Determine k_n
 - Effect of h_n (window width): A compromised value for a fixed number of training examples should be chosen

$$p_n(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{\mathbf{x} - \mathbf{x}_i}{h_n}\right) \quad (V_n = h_n^d)$$

$\varphi(\cdot)$ being a pdf function  $p_n(\cdot)$ being a pdf function

window function
(being pdf) $\varphi(\cdot)$ + window width h_n + training data \mathbf{x}_i  Parzen pdf $p_n(\cdot)$

Summary (Cont.)

- k_n -nearest-neighbor: Fix $k_n \rightarrow$ Determine V_n

specify $k_n \rightarrow$ center a cell about $x \rightarrow$ grow the cell until capturing k_n nearest examples \rightarrow return cell volume as V_n

The principled rule to specify k_n [pp.175]

$$\lim_{n \rightarrow \infty} k_n = \infty$$

$$\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$$



A rule-of-thumb
choice for k_n :

$$k_n = \sqrt{n}$$

Summary (Cont.)

■ Nearest neighbor (NN) rule & distance metric

- Classification with NN rule: Voronoi tessellation
- Error bounds of NN rule w.r.t. Bayes risk

$$P^*(e) \leq P(e) \leq P^*(e) \left(2 - \frac{c}{c-1} P^*(e) \right)$$

- Classification with k NN rule
- Reducing computational complexity
 - *Partial distance, pre-structuring, Editing/Pruning/Condensing*
- Distance metric
 - *non-negativity, reflexivity, symmetry, triangle inequality*
 - *Minkowski distance, Tanimoto distance, Hausdorff distance*