

- Chapter 4: Nonparametric Techniques
 - 4.1 Introduction
 - 4.2 Probability Density Estimation
 - 4.3 Parzen Windows
 - 4.4 k_n -nearest-neighbor
 - 4.5 Nearest-Neighbor Rule
 - 4.5.1 Definition
 - 4.4.2 Voronoi tessellation (维诺图)
 - 4.4.3 Error Rate of Nearest Neighbor Rule
 - 4.4.4 Error Bounds of Nearest Neighbor Rule
 - 4.4.5 k_n -nearest-neighbor Rule
 - 4.4.6 Computational Complexity of KNN Rule
 - 4.5 Distance Metric
 - 4.5.1 Properties of Distance Metric
 - 4.5.2 Euclidean Distance
 - 4.5.3 Distance Metric Between Sets

Chapter 4: Nonparametric Techniques

4.1 Introduction

Potential problems for parametric techniques:

- The assumed parametric form **may not fit the ground-truth density** encountered in practice, e.g:
 - **Assumed parametric form:** Unimodal(such as Gaussian pdf)
 - **Ground-truth density:** Multimodal(such as mixture of Gaussian pdfs)

So here comes the nonparametric techniques: $p(x|w_j)$ doesn't have parametric form. We use the data to speak for themselves.

1. **Parzen Windows**
2. **k_n -nearest-neighbor**

4.2 Probability Density Estimation

General Settings

- **Feature space:** $\mathcal{F} = \mathbf{R}^d$
- **Feature vector:** $x \in \mathcal{F}$
- **pdf function:** $x \sim p(\cdot)$

Fundamental Fact: The probability (**a smoothed/averaged version of $p(x)$**) of a vector x falling into a region $\mathcal{R} \subset \mathcal{F}$:

$$P = \Pr[x \in \mathcal{R}] = \int_{\mathcal{R}} p(x') dx'$$

Given n examples (i.i.d.)

$$x_1, x_2, \dots, x_n$$

with $x_i \sim p(\cdot)$ ($1 \leq i \leq n$)


Let X be the (discrete) **random variable** representing the number of examples falling into \mathcal{R} , so X will take **Binomial Distribution**: $X \sim \mathcal{B}(n, P)$ and we have:

$$\begin{aligned} \Pr[X = r] &= \binom{n}{r} P^r (1 - P)^{n-r} \\ \mathbb{E}[X] &= nP \quad \Rightarrow \quad P = \frac{\mathbb{E}[X]}{n} \end{aligned}$$

If we assume that $p(x)$ is constant and region \mathcal{R} is so small that $p(\cdot)$ hardly varies with \mathcal{R} , then we have:

$$P = \Pr(x \in \mathcal{R}) = \int_{\mathcal{R}} p(x') dx' \approx p(x) \int_{\mathcal{R}} 1 dx' = p(x)V$$

- x is a point in \mathcal{R}
- V is the volume enclosed by \mathcal{R}



- k/n : success / total
- all curve converge to a **single peak** at $\mathbb{E}(X)$
- when n get larger, curve gets sharper around $\mathbb{E}(X)$

X peaks sharply about $\mathbb{E}(X)$ when n is large enough.

Let k be the actual value of X after observing the i.i.d. training examples x_1, x_2, \dots, x_n

Given those evidence, we have $k \approx \mathbb{E}[X]$, hence we can conclude that:

$$p(x) = \frac{k/n}{V}$$

If we fixed the value of V and are able to get more training examples, then the value of k/n will converge to the true value of $p(x)$, which is a **smoothed version**:

$$\frac{P}{V} = \frac{\int_{\mathcal{R}} p(x') dx'}{\int_{\mathcal{R}} dx'}$$

However, we hope to estimate $p(x)$ which is not a smoothed version, so we need to make V approach 0, which lead to the following problems:

- $V \rightarrow 0$, the region \mathcal{R} will become so small that there may be no training examples falling into the region, which makes no sense.
- If 1 or 2 sample accidentally fall into the region \mathcal{R} , then the estimation will become infinite, which is also not reasonable.

Theoretically, we construct a series of regions $\mathcal{R}_1, \mathcal{R}_2, \dots, \mathcal{R}_m$ which contain the dot x to overcome the limitations. The first region use 1 sample, the second region use 2 samples, and so on. Hence we have:

$$p_n(x) = \frac{k_n/n}{V_n}$$

- V_n : the volume of region \mathcal{R}_n
- n : the number of training examples
- k_n : the number of training examples falling into \mathcal{R}_n

Requirements to be satisfied, which make $p_n(x)$ converge to $p(x)$:

1. $\lim_{n \rightarrow \infty} V_n = 0$: ensure the space-smoothed P/V converge to $p(x)$ under the situation of uniform convergence in the region and continuity of $p(\cdot)$ at point x .
2. $\lim_{n \rightarrow \infty} k_n = \infty$: meaningful only when $p(x) \neq 0$ and guarantees that the ratio of frequencies can converge to the probability P
3. $\lim_{n \rightarrow \infty} k_n/n = 0$

Methods:

1. **Parzen Windows**: Fix V_n and determine k_n
2. **k_n -nearest-neighbor**: Fix k_n and determine V_n

4.3 Parzen Windows

Assumptions:

- \mathcal{R}_n : a d-dimensional hypercube
- h_n : the length of the edge of \mathcal{R}_n

Determine k_n with **window function**, a.k.a. **kernel function**, **potential function**, etc.

- **Properties of window functions**
 - non-negativity: $\varphi(\mathbf{u}) \geq 0$
 - uniformity: $\int \varphi(\mathbf{u}) d\mathbf{u} = 1$
- **Extension**
 - $\varphi(\cdot)$ is not limited to be the following hypercube window function but any pdf function that satisfies the above two properties.

Define window function:

$$\varphi(\mathbf{u}) = \begin{cases} 1 & |u_j| \leq \frac{1}{2}, j = 1, \dots, d, \\ 0 & \text{otherwise} \end{cases}$$

$\varphi(\mathbf{u})$ defines a **unit hypercube** centered at the origin and $V_n = h_n^d$

$\varphi\left(\frac{x-x_i}{h_n}\right) = 1 \Leftrightarrow x_i$ falls within the hypercube of volume V_n centered at x with edge length h_n

- $x - x_i$: the relative position between the sample point x_i and the target point x
- h_n : bandwidth parameter which controls the width of the window

Hence the total number of the sample points falling into the hypercube can be calculated as:

$$k_n(x) = \sum_{i=1}^n \varphi\left(\frac{x-x_i}{h_n}\right)$$

According to the probability density, we have:

$$p_n(x) = \frac{k_n/n}{V_n} = \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{x-x_i}{h_n}\right)$$

- **meaning**: the average contribution of the sample points to the target point x
- The kernel function φ weights the influence of each sample point x_i on x :
 1. If x_i is very close to x (i.e., $|x - x_i|$ is small), then $\varphi\left(\frac{x-x_i}{h_n}\right)$ will approach 1.
 2. If x_i is far from x , $\varphi\left(\frac{x-x_i}{h_n}\right)$ will approach 0.
 3. $\varphi(\cdot)$ is a pdf function $\Rightarrow p_n(x)$ is a pdf function.

The proof of 3 is as follows:

- **Notice**: x is a multi-dimensional vector and the window bandwidth h_d do the same for all the dimensions. When using the substitution $\mathbf{u} = \frac{x-x_i}{h_n}$, we get $dx = h_n^d d\mathbf{u}$

$$\begin{aligned} \varphi(\cdot) \geq 0 \Rightarrow p_n(x) &= \frac{1}{n} \sum_{i=1}^n \frac{1}{V_n} \varphi\left(\frac{x-x_i}{h_n}\right) \geq 0 \\ \int p_n(x) dx &= \frac{1}{nV_n} \sum_{i=1}^n \int \varphi\left(\frac{x-x_i}{h_n}\right) dx \\ &= \frac{1}{nV_n} \sum_{i=1}^n \int h_n^d \varphi\left(\frac{x-x_i}{h_n}\right) d\left(\frac{x-x_i}{h_n}\right) \\ &= \frac{V_n}{nV_n} \sum_{i=1}^n \int \varphi(\mathbf{u}) d\mathbf{u} = 1 \end{aligned}$$

Above all we can conclude that:

- **windows function (being pdf)** $\varphi(\cdot)$ + **window width** h_n + **training data** x_i = **Parzen pdf** $p_n(\cdot)$

Now we exploit the effect of window width h_n on the Parzen pdf $p_n(x)$. We define function $\delta_n(x)$ as:

$$\delta_n(x) = \frac{1}{V_n} \varphi\left(\frac{x}{h_n}\right) = \frac{1}{h_n^d} \varphi\left(\frac{x}{h_n}\right)$$

- window width h_n affect two key attributes of partial contribution function $\delta_n(\cdot)$:
 1. **Amplitude**: determined by $\frac{1}{h_n^d}$, the larger h_n , the smaller the amplitude.
 2. **Width**: determined by $\frac{x}{h_n}$, the larger h_n , the wider the width. (**larger width cause slower change so as to increase effective range**)
- $\delta_n(\cdot)$ being a pdf function

$$\begin{aligned}\delta_n(\cdot) &\geq 0 \\ \int \delta_n(x) dx &= \int \frac{1}{h_n^d} \varphi\left(\frac{x}{h_n}\right) dx \\ &= \int \frac{1}{h_n^d} \cdot \varphi(\mathbf{u}) \cdot h_n^d d\mathbf{u} = 1\end{aligned}$$

We can rewrite $p_n(x)$ as:


$$p_n(x) = \frac{1}{n} \sum_{i=1}^n \delta_n(x - x_i)$$

- $p_n(x)$: **superposition** of n sample contributions interpolations
- x_i : contributes to $p_n(x)$ based on the **distance** between x and x_i
- **Influence**:
 - h_n **very large**
 1. $\delta_n(x)$ being broad with small amplitude
 2. $p_n(x)$ will be the superposition of n broad, slowly changing functions, i.e. **being smooth with low resolution**
 3. may cause **underfitting** and be unable to capture local structure
 - h_n **very small**
 1. $\delta_n(x)$ being sharp with large amplitude
 2. $p_n(x)$ will be the superposition of n sharp pulses, i.e. **being variable/unstable with high resolution**
 3. may cause **overfitting** and can be more affected by noise

A compromised value of h_n should be sought for limited number of training examples.

Suppose $\varphi(\cdot)$ being a 2-d Gaussian pdf and $n=5$


The shape of $p_n(x)$ with decreasing values of h_n



4.4 k_n -nearest-neighbor

$$p_n(x) = \frac{k_n/n}{V_n}$$

- Fix k_n then determine V_n
- **Procedure:** specify $k_n \rightarrow$ center a cell about $x \rightarrow$ grow the cell until capturing k_n nearest examples \rightarrow return cell volumes as V_n
- The principled rule to specify k_n (the **sufficient and necessary conditions** of $p_n(x)$ converging to $p(x)$):
 1. $\lim_{n \rightarrow \infty} k_n = \infty$
 2. $\lim_{n \rightarrow \infty} \frac{k_n}{n} = 0$
- A rule-of-thumb (经验法则) choice for k_n : $k_n = \sqrt{n}$



- 1-dimension example
 - Smaller k_n produces detailed but noisier estimates
 - Larger k_n creates smoother estimates but with reduced local detail.
- 2-dimension example
 - The $k_n = 5$ estimation dynamically adapts to the data density, showing higher values in dense areas and lower values in sparse regions.

4.5 Nearest-Neighbor Rule

4.5.1 Definition


Given:

- The **label space**: $\Omega = \{\omega_1, \omega_2, \dots, \omega_c\}$
- A set of n labeled training examples: $\mathcal{D}^n = \{(\mathbf{x}_i, \theta_i) \mid 1 \leq i \leq n\}$, where:
 - $\mathbf{x}_i \in \mathbb{R}^d$: The i -th feature vector in d -dimensional space
 - $\theta_i \in \Omega$: The label associated with \mathbf{x}_i

For a test example \mathbf{x} :

1. Identify its nearest neighbor \mathbf{x}' : $\mathbf{x}' = \arg \min_{\mathbf{x}_i \in \{\mathbf{x}_1, \dots, \mathbf{x}_n\}} D(\mathbf{x}_i, \mathbf{x})$,
 - $D(\mathbf{a}, \mathbf{b})$: distance metric between two vectors \mathbf{a} and \mathbf{b} , such as the Euclidean distance.
2. Assign the label θ' associated with \mathbf{x}' to \mathbf{x} .

4.4.2 Voronoi tessellation (维诺图)



Each training example \mathbf{x} leads to a cell in the Voronoi tessellation

- any point in the cell is closer to \mathbf{x} than to any other training examples
- partition the feature space into n cells
- any point in the cell share the same class label as \mathbf{x}

4.4.3 Error Rate of Nearest Neighbor Rule

$$P(e) = \int P(e|x)p(x)dx$$

- $P(e)$: The **average probability of error** based on nearest-neighbor rule
- $P(e | x)$: The probability of making an erroneous classification on x based on nearest-neighbor rule
- $p(x)$: The probability density function of the feature space, which describes how data points are distributed

According to the **Bayes Decision Rule**, we can minimize the error rate at each point x by choosing the class label w_m with the maximum posterior probability:

$$P^*(e|x) = 1 - P(w_m|x)$$

- $P(w_m|x)$: the posterior probability of the most likely class at x , i.e. $P(w_m|x) = \max_{1 \leq i \leq c} P(w_i|x)$

- $P^*(e|x)$: the **minimum** possible value of $P(e|x)$, i.e. the one given by Bayesian decision rule

The Bayes risk(under zero-one loss) can be written as:

$$P^*(e) = \int P^*(e|x)p(x)dx$$

4.4.4 Error Bounds of Nearest Neighbor Rule

$$P^*(e) \leq P(e) \leq P^*(e)\left(2 - \frac{c}{c-1}P^*(e)\right)$$

- $P_n(e)$: The probability of error under the NN rule with n training examples
- $P = \lim_{n \rightarrow \infty} P_n(e)$: the asymptotic probability of error as the number of training examples approaches infinity
- c : the number of class labels

A conditional probability of error $P(e|x, x')$ relying on testing sample x and the nearest neighbor x' :

$$P(e|x) = \int P(e|x, x')p(x'|x)dx'$$

When $\theta \neq \theta'_n$, it produce a classification error:

$$\begin{aligned} P_n(e|x, x'_n) &= 1 - \sum_{i=1}^c P(\theta = w_i, \theta'_n = w_i | x, x'_n) \\ &= 1 - \sum_{i=1}^c P(w_i|x)P(w_i|x'_n) \end{aligned}$$

According to the fact that $p(x'|x)$ peaks sharp around x but stay low at other locations, we can assume $p(x'|x)$ approach a **Dirac delta function** δ centered at x when n approaches infinity.

Hence we have:

- When $\delta(x - x_0)$ being a integral kernel, we have: $\int f(x)\delta(x - x_0)dx = f(x_0)$
- When $n \rightarrow \infty$, $P(w_i|x'_n) \rightarrow P(w_i|x)$, hence we have $\sum_{i=1}^c P(w_i|x)P(w_i|x'_n) \rightarrow \sum_{i=1}^c P^2(w_i|x)$

$$\begin{aligned} \lim_{n \rightarrow \infty} P_n(e|x, x'_n) &= \int \left[1 - \sum_{i=1}^c P(w_i|x)P(w_i|x'_n)\right] \delta(x' - x) dx' \\ &= 1 - \sum_{i=1}^c P^2(w_i|x) \end{aligned}$$

Hence we have asymptotic probability of NN error:

$$\begin{aligned}
P &= \lim_{n \rightarrow \infty} P_n(x) \\
&= \lim_{n \rightarrow \infty} \int P_n(e|x)p(x)dx \\
&= \int \left[1 - \sum_{i=1}^c P^2(w_i|x) \right] p(x)dx
\end{aligned}$$

Proof of the lower bound:

$P^*(e)$ is the minimum possible value of $P(e)$ given by the Bayesian decision rule, which is the optimal classifier.

Proof of the upper bound:

$$\sum_{i=1}^c P^2(w_i|x) = P^2(w_m|x) + \sum_{i \neq m} P^2(w_i|x)$$

- **Constraints:**

- $P(w_i|x) \geq 0$
- $\sum_{i \neq m} P(w_i|x) = 1 - P(w_m|x) = P^*(e|x)$

To make the minimum of $\sum_{i=1}^c P^2(w_i|x)$, we let:

$$P(w_i|x) = \begin{cases} \frac{P^*(e|x)}{c-1} & i \neq m \\ 1 - P^*(e|x) & i = m \end{cases}$$

Hence we have :

$$\begin{aligned}
\sum_{i=1}^c P^2(w_i|x) &\geq (1 - P^*(e|x))^2 + \frac{P^{*2}(e|x)}{c-1} \\
\therefore 1 - \sum_{i=1}^c P^2(w_i|x) &\leq 2P^*(e|x) - \frac{c}{c-1} P^{*2}(e|x)
\end{aligned}$$

Directly we can get: $P \leq 2P^*$

Notice that:

$$\begin{aligned}
Var[P^*(e|x)] &= E(P^{*2}(e|x)) - E(P^*(e|x))^2 \\
&= \int P^{*2}(e|x)p(x)dx - (\int P^*(e|x)p(x)dx)^2 \\
&= \int P^{*2}(e|x)p(x)dx - P^{*2} \geq 0
\end{aligned}$$

therefore we have:

$$\int P^{*2}(e|x)p(x)dx \geq P^{*2}$$

According to the fact we get before:


$$\begin{aligned}
P &= \int \left[1 - \sum_{i=1}^c P^2(w_i|x) \right] p(x) dx \\
&\leq \int \left[2P^*(e|x) - \frac{c}{c-1} P^{*2}(e|x) \right] p(x) dx \\
&= 2P^* - \int \frac{c}{c-1} P^{*2}(e|x) p(x) dx \\
&\leq 2P^* - \frac{c}{c-1} P^{*2} = P^*(2 - \frac{c}{c-1} P^*)
\end{aligned}$$

4.4.5 k_n -nearest-neighbor Rule

Classification with k-nearest-neighbor rule

- Given the label space $\Omega = \omega_1, \omega_2, \dots, \omega_c$ and a set of n labeled training examples $D^n = (x_i, \theta_i) \quad 1 \leq i \leq n$, where $x_i \in \mathbb{R}^d$ and $\theta_i \in \Omega$:
 - For test example x , identify $S' = \{x_i \mid x_i \text{ is among the } KNN \text{ of } x\}$ and then assign the most frequent label w.r.t. S' , i.e., $\arg \max_{\omega_i \in \Omega} \sum_{x_i \in S'} 1_{\theta_i = \omega_i}$ to x .
 - 1_θ : an indicator function which returns 1 if predicate θ right, and 0 otherwise.

For binary classification problem ($c = 2$), and odd value of k is generally used to avoid ties.



4.4.6 Computational Complexity of KNN Rule

Given n labeled training examples in d -dimensional space, the computational complexity of **classifying one test example** is $O(dn)$

General ways of reducing computational burden:

- Calculate Partial Distance:** $D_r(a, b) = \left(\sum_{j=1}^r (a_j - b_j)^2 \right)^{\frac{1}{2}} \quad (r < d)$
 - Using a subset r of the total **d-dimensional space**
 - As more dimensions are gradually added, the values of partial distances are strictly **non-decreasing**
- Pre-structuring:** create some form of search tree, where nearest neighbors are **recursively** identified following the tree structure (each root node corresponds to a branch of the

discriminant region)

- This may cause error because when the testing sample is near the boundary of the region, the nearest neighbor may not be found in the discriminated region.
- **Editing/Pruning/Condensing:** eliminate "redundant" ("useless") examples from the training set, e.g. example surrounded by training examples of the same class label

4.5 Distance Metric

4.5.1 Properties of Distance Metric

The NN/kNN rule depends on the use of distance metric to identify nearest neighbors.

Four properties of distance metric:


- **Non-negativity:** $D(a, b) \geq 0$
- **Reflexivity (自反性):** $D(a, b) = 0 \Leftrightarrow a = b$
- **Symmetry:** $D(a, b) = D(b, a)$
- **Triangle Inequality:** $D(a, b) + D(b, c) \geq D(a, c)$

4.5.2 Euclidean Distance

Euclidean distance:

$$D(a, b) = \left(\sum_{j=1}^d (a_j - b_j)^2 \right)^{\frac{1}{2}}$$

Potential Issue:



Scaling the features → change the distance relationship

- **Issue:** If the coordinate system is transformed, the new distance relationship calculated by Euclidean distance may be very different from the original one.
- **Possible solution:** normalize each feature into equal-sized intervals, e.g. [0,1]

Minkowski Distance Metric (a.k.a. L_k norm)

$$L_k(a, b) = \left(\sum_{j=1}^d |a_j - b_j|^k \right)^{\frac{1}{k}} \quad (k > 0)$$

- $k = 2$: Euclidean Distance
- $k = 1$: Manhattan Distance (city block distance)
- $k = \infty$: L_∞ distance

- $L_\infty(a, b) = \max_{1 \leq j \leq d} |a_j - b_j|$

4.5.3 Distance Metric Between Sets

Tanimoto distance:

$$D_{\text{Tanimoto}}(S_1, S_2) = \frac{n_1 + n_2 - 2n_{12}}{n_1 + n_2 - n_{12}}$$

$$(n_1 = |S_1|, n_2 = |S_2|, n_{12} = |S_1 \cap S_2|)$$

Hausdorff Distance:

$$D_H(S_1, S_2) = \max \left(\max_{s_1 \in S_1} \min_{s_2 \in S_2} D(s_1, s_2), \max_{s_2 \in S_2} \min_{s_1 \in S_1} D(s_2, s_1) \right)$$

$$(D(s_1, s_2) : \text{any distance metric between } s_1 \text{ and } s_2)$$

- $\max_{s_1 \in S_1} \min_{s_2 \in S_2} D(s_1, s_2)$: for every s_1 in S_1 , find the maximum value of the minimum distance between s_1 and all s_2 in S_2

Example: Hausdorff Distance between two sets of feature vectors

$$S_1 = \{(0.1, 0.2)^t, (0.3, 0.8)^t\} \quad S_2 = \{(0.5, 0.5)^t, (0.7, 0.3)^t\}$$

- $D_H(S_1, S_2) = \max \left(\max_{s_1 \in S_1} \min_{s_2 \in S_2} D(s_1, s_2), \max_{s_2 \in S_2} \min_{s_1 \in S_1} D(s_2, s_1) \right)$
 - $\max_{s_1 \in S_1} \min_{s_2 \in S_2} D(s_1, s_2)$
 - for $s_1 = (0.1, 0.2)^t$
 - $d((0.1, 0.2)^t, (0.5, 0.5)^t) = 0.5$
 - $d((0.1, 0.2)^t, (0.7, 0.3)^t) = \sqrt{0.37}$
 - for $s_1 = (0.3, 0.8)^t$
 - $d((0.3, 0.8)^t, (0.5, 0.5)^t) = \sqrt{0.13}$
 - $d((0.3, 0.8)^t, (0.7, 0.3)^t) = \sqrt{0.41}$
 - $\max_{s_1 \in S_1} \min_{s_2 \in S_2} D(s_1, s_2) = \max \left(\min(0.5, \sqrt{0.37}), \min(\sqrt{0.13}, \sqrt{0.41}) \right) = \max(0.5, \sqrt{0.13}) = 0.5$
 - $\max_{s_2 \in S_2} \min_{s_1 \in S_1} D(s_2, s_1)$
 - for $s_2 = (0.5, 0.5)^t$
 - $d((0.5, 0.5)^t, (0.1, 0.2)^t) = 0.5$
 - $d((0.5, 0.5)^t, (0.3, 0.8)^t) = \sqrt{0.13}$
 - for $s_2 = (0.7, 0.3)^t$
 - $d((0.7, 0.3)^t, (0.1, 0.2)^t) = \sqrt{0.37}$
 - $d((0.7, 0.3)^t, (0.3, 0.8)^t) = \sqrt{0.41}$
 - $\max_{s_2 \in S_2} \min_{s_1 \in S_1} D(s_2, s_1) = \max \left(\min(0.5, \sqrt{0.13}), \min(\sqrt{0.37}, \sqrt{0.41}) \right) = \max(\sqrt{0.13}, \sqrt{0.37}) = \sqrt{0.37}$
- $D_H(S_1, S_2) = \max(\sqrt{0.5}, \sqrt{0.37}) = \sqrt{0.37} \approx 0.61$